



(a) Model ladder: one refinement per rung



*each rung trains a CA arm and a
parameter-matched TC arm identically*

(b) Donor-level architecture (R3 shown)



attention weights are diagnostics, not explanations